

Lexical Analysis

What is Lexical Analysis

- ❑ Comes from the word lexicon, or dictionary
- ❑ **Lexical Analysis:** Partitioning a string into words
 - These words are also called **tokens**
- ❑ We will use this code as a running example:

```
if (i==j)
    z = 0;
else
    z = 1;
```
- ❑ Code is provided as input string to lexical analysis
 - `"if(i == j)\n\tz = 0; \nelse\n\tz = 1; \n"`
- ❑ Goal is turning the string into tokens, or tokenization

What is a Token ?

- ❑ Smallest unit that has meaning in a string
 - In English, tokens are English words:
nouns, verbs, adjectives, ...
 - In a programming language:
identifiers, integers, keywords, whitespace, ...

- ❑ A token is a tuple (`type`, `lexeme`)
 - `type`: the token type that the token belongs to
 - Identifier: string of letters and digits, starting with a letter
 - Integer: string of digits
 - Keyword: “else”, “if”, “while”, ...
 - Whitespace: string of blanks, newlines, and tabs
 - `lexeme`: actual string value of this token

Lexical Analysis is the act of Tokenization

- ❑ Output of lexical analysis is a stream of tokens
- ❑ Tokens are the input to Syntax Analysis (a.k.a. Parsing)
 - Parsers rely on token type to figure out role of each token
E.g. a keyword is treated differently from an identifier

Lexical Analysis Tokenization Example

❑ Given `"if(i == j)\n\tz = 0; \tnelse\n\tz = 1; \n"`

❑ What would be an output of lexical analysis?

Recall a token is a tuple `(type, lexeme)`

❑ Output:

`(keyword, "if")``(left-parenthesis, "(")``(identifier, "i")``(equals-op, "=")``(identifier, "j")``(right-parenthesis, ")")``(whitespace, "\n\t")``(identifier, "z")``(assign-op, "=")``(integer, "0")``(semicolon, ";")``(whitespace, "\n")``(keyword, "else")``(whitespace, "\n\t")``(identifier, "z")``(assign-op, "=")``(integer, "1")``(semicolon, ";")``(whitespace, "\n")`

❑ The lexer usually discards “non-interesting” tokens that don’t contribute to parsing, e.g., whitespace, comments

Some language features makes lexing difficult

- ❑ FORTRAN compilation rule: **whitespace is insignificant**
 - Reason: inaccuracy of card punching by operators
- ❑ Consider
 - DO 5I=1,25
 - DO 5I=1.25
- ❑ This is the interpretation of the two statements:
 - Former: Iterate from I=1 to I=25 with step size 5
 - Latter: Assign 1.25 to variable DO5I
- ❑ Reading left-to-right, cannot tell if DO5I is a variable or DO statement; Have to continue until “,” or “.” is reached.
 - “lookahead” may be required to decide on tokens
 - Feedback necessary from parser to lexical analysis

C++ language has difficult features too

- ❑ C++ has the Right Angle Brackets >> issue:
<https://www.open-std.org/jtc1/sc22/wg21/docs/papers/2005/n1757.html>
- ❑ `typedef std::vector<std::vector<int> > Table; // OK`
`typedef std::vector<std::vector<int>> Table; // Error`
- ❑ Why? >> is read as one token which can either be:
 - A stream operator (e.g. `cin >> var`)
 - Or a shift operator (e.g. `1 >> 2`)
- ❑ Space is needed in between to create two > tokens
- ❑ Fixed in C+11 standard so this is no longer an error
 - That makes tokenization decision on >> context dependent
 - Again forcing lexical analysis to get feedback from parser

Lexical Analysis Implementation

Step 1:

- Define a set of token types
 - Refer to language specifications
 - Types you choose depends on design of parser
 - Recall “ $if(i == j) \backslash n \backslash tz = 0; \backslash n else \backslash n \backslash tz = 1; \backslash n$ ”
 - Should “==” be one token? or two tokens? Depends.
 - Should “if”, “then”, “else” be separate types or just one keyword type? Depends.

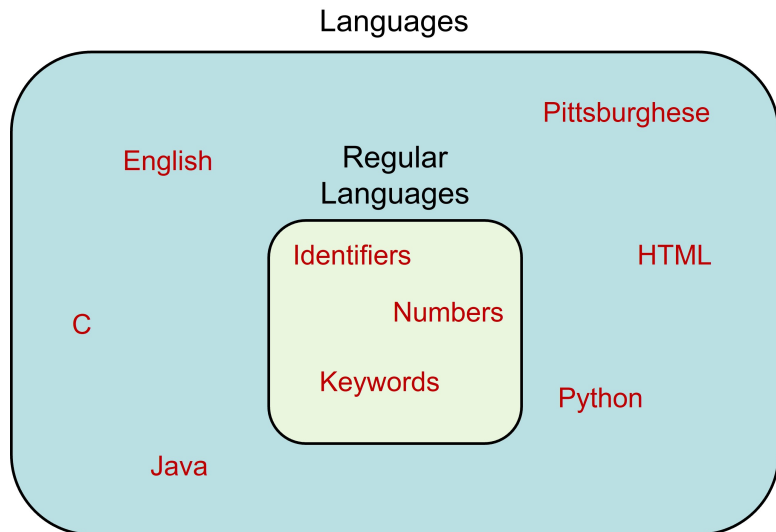
Step 2:

- For each token type, describe which string belongs to it

Describing strings belonging to a token type

- ❑ A token type is something that looks like this:
Identifier: string of letters and digits, starting with a letter
- ❑ Is there a more formal (mathematical) way to express this?
 - Yes! By using the formalism of **Languages**.
- ❑ Definition of **Language**:
Let Σ be a set of characters, a **language** over Σ is a set of strings of the characters drawn from Σ
 - So by definition, any token type is a Language
 - And so are English, Java, Python, and HTML
- ❑ Some Languages can be very difficult to express formally
 - Imagine having to formally describe the English language!

Token types belong to a subset of Languages (called Regular Languages)



Regular Expressions express Regular Languages

□ Definition of **Regular Expression**

The **regular expressions (REs)** over Σ are the total set of expressions that can be constructed using the following components:

- ϵ
- 'c' where $c \in \Sigma$
- $A + B$ where A, B are **RE** over Σ
- AB where A, B are **RE** over Σ
- A^* where A is a **RE** over Σ

□ **Regular Languages** are defined as languages that can be expressed using **Regular Expressions**.

Atomic Regular Expressions

- Single character denotes a set of one string
 $'c' = \{ "c" \}$
- Epsilon* or ϵ character denotes a zero length string $\epsilon = \{ "" \}$
- Empty set is $\{ \} = \phi$, not the same as ϵ
 $\text{size}(\phi) = 0$
 $\text{size}(\epsilon) = 1$
 $\text{length}(\epsilon) = 0$

Compound Regular Expressions

- Union: if A and B are REs, then

$$A + B = \{ s \mid s \in A \text{ or } s \in B \}$$

- Concatenation of sets/strings

$$AB = \{ ab \mid a \in A \text{ and } b \in B \}$$

- Iteration (Kleene closure)

$$A^* = \bigcup_{i \geq 0} A^i \quad \text{where } A^i = A \dots A \text{ (} i \text{ times)}$$

in particular

$$A^* = \{ \epsilon \} + A + AA + AAA + \dots$$

$$A^+ = A + AA + AAA + \dots = A A^*$$


The $L(\textit{Expression})$ Notation

□ $L(\textit{Expression})$: means the language defined by *Expression*

□ Some example languages defined using the $L()$ notation:

- $L(\epsilon) = \{ \epsilon \}$
- $L('c') = \{ c \}$
- $L(A+B) = L(A) \cup L(B)$
- $L(AB) = \{ ab \mid a \in L(A) \text{ and } b \in L(B) \}$
- $L(A^*) = \bigcup_{i \geq 0} L(A^i)$

Examples

 Keywords: “else” or “if” or “while” or ...

Examples

Keywords: “else” or “if” or “while” or ...

➤ ‘else’ + ‘if’ + ‘while’ + ...

➤ ‘else’ abbreviates

‘e’ (concatenate) ‘l’ (concatenate) ‘s’ (concatenate) ‘e’

Examples

Keywords: “else” or “if” or “while” or ...

➤ ‘else’ + ‘if’ + ‘while’ + ...

➤ ‘else’ abbreviates

‘e’ (concatenate) ‘l’ (concatenate) ‘s’ (concatenate) ‘e’

➤ keywords = { ‘else’, ‘if’, ‘then’, ‘while’, ... }

Examples

Keywords: “else” or “if” or “while” or ...

- ‘else’ + ‘if’ + ‘while’ + ...
- ‘else’ abbreviates
‘e’ (concatenate) ‘l’ (concatenate) ‘s’ (concatenate) ‘e’
- keywords = { ‘else’, ‘if’, ‘then’, ‘while’, ... }

Integer

- digit = ‘0’ + ‘1’ + ‘2’ + ‘3’ + ‘4’ + ‘5’ + ‘6’ + ‘7’ + ‘8’ + ‘9’
- integer = digit digit*

Examples

Keywords: “else” or “if” or “while” or ...

- ‘else’ + ‘if’ + ‘while’ + ...
- ‘else’ abbreviates
‘e’ (concatenate) ‘l’ (concatenate) ‘s’ (concatenate) ‘e’
- keywords = { ‘else’, ‘if’, ‘then’, ‘while’, ... }

Integer

- digit = ‘0’ + ‘1’ + ‘2’ + ‘3’ + ‘4’ + ‘5’ + ‘6’ + ‘7’ + ‘8’ + ‘9’
- integer = digit digit*
 - **Q:** is ‘000’ an integer?

Examples

Keywords: “else” or “if” or “while” or ...

- ‘else’ + ‘if’ + ‘while’ + ...
- ‘else’ abbreviates
‘e’ (concatenate) ‘l’ (concatenate) ‘s’ (concatenate) ‘e’
- keywords = { ‘else’, ‘if’, ‘then’, ‘while’, ... }

Integer

- digit = ‘0’ + ‘1’ + ‘2’ + ‘3’ + ‘4’ + ‘5’ + ‘6’ + ‘7’ + ‘8’ + ‘9’
- integer = digit digit*
 - **Q:** is ‘000’ an integer?
 - **Q:** how to define another integer RE that excludes sequences with leading 0s?

More Examples

- ❑ Identifier: strings of letters or digits, starting with a letter
 - letter = 'A' + ... + 'Z' + 'a' + ... + 'z'
 - Identifier = letter (letter + digit)*

- ❑ Whitespace: a non-empty string of blanks, newlines, tabs
 - whitespace = (' ' + '\n' + '\t') +

More Examples

❏ Identifier: strings of letters or digits, starting with a letter

- letter = 'A' + ... + 'Z' + 'a' + ... + 'z'
- Identifier = letter (letter + digit)*
 - **Q:** is letter (letter* + digit*) the same?

❏ Whitespace: a non-empty string of blanks, newlines, tabs

- whitespace = (' ' + '\n' + '\t') +

More Examples

☐ Phones number: consider (412) 624-0000

- $\Sigma = \text{digit} \cup \{ -, (,) \}$
- $\text{area} = \text{digit}^3$
- $\text{exchange} = \text{digit}^3$
- $\text{phone} = \text{digit}^4$
- $\text{phoneNumber} = '(' \text{ area } ')' \text{ exchange } '-' \text{ phone}$

☐ Email address: student @ pitt.edu

- $\Sigma = \text{letter} \cup \{ ., @ \}$
- $\text{name} = \text{letter}^+$
- $\text{emailAddress} = \text{name } '@' \text{ name } '.' \text{ name}$

More Regular Expression Notations

□ Some “syntactic sugar” for regular expressions:

Union: $A + B \equiv A \mid B$

Option: $A + \varepsilon \equiv A ?$

Range: $'a' + 'b' + \dots + 'z' \equiv [a-z]$

Excluded range:
complement of $[a-z] \equiv [^\wedge a-z]$

Implementation of Lexical Analysis

- We have learnt the formalism for lexical analysis
 - Regular expression (RE)

Implementation of Lexical Analysis

- We have learnt the formalism for lexical analysis
 - Regular expression (RE)
- How do we get the actual lexical analyzer?

Implementation of Lexical Analysis

- ❑ We have learnt the formalism for lexical analysis
 - Regular expression (RE)

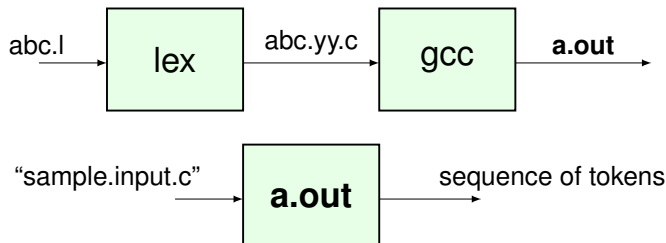
- ❑ How do we get the actual lexical analyzer?
 - **Solution 1:** Convert RE to code using a tool — Lex (for C), Flex (for C++), Jlex (for java)
 - Programmer specifies tokens of interest using REs
 - The tool generates the source code from the given REs

Implementation of Lexical Analysis

- ❑ We have learnt the formalism for lexical analysis
 - Regular expression (RE)

- ❑ How do we get the actual lexical analyzer?
 - **Solution 1:** Convert RE to code using a tool — Lex (for C), Flex (for C++), Jlex (for java)
 - Programmer specifies tokens of interest using REs
 - The tool generates the source code from the given REs
 - **Solution 2:** Write the code manually from REs
 - This is likely similar to code that the tool generates
 - Table-driven code based on a finite state automaton

Lex: a Tool for Lexical Analysis



- ❑ Big difference from your previous coding experience
 - Writing REs instead of the code itself
 - Writing actions associated with each RE
- ❑ We will first describe structure of specification file
- ❑ The internals of the tool will be discussed later

Example Lex Specification File

```
/* 1. Regular expression definitions section */
%{
/* Code block inserted for includes and declarations */
#include <stdlib.h>
}%
string    [a-z]+
space     [ ]+
%%

/* 2. Rules section: action for each regular expression */
{string}  { printf("lexeme: %s, len=%d\n", yytext, yyleng); }
{space}   { /* No action */ }
%%

/* 3. User code section */
int main() {
    while( yylex() != 0 ) {}
    return 0;
}
```

Example Lex Specification File: Explanation

Overview of operation:

- 1 Parser calls `yylex()` when ready to process the next token
- 2 `yylex()` tokenizes longest string that matches an RE
- 3 `yylex()` stores the token lexeme in `yytext`
- 4 `yylex()` stores length of lexeme in `yyleng`
- 5 `yylex()` executes the action `{ ... }` in the rule for RE
- 6 `yylex()` returns 0 if no more tokens / non-zero otherwise

To test the lexer without a parser, we need a lexer driver

- The `main()` function serves as the lexer driver
- On piping "hello world!" to input of lexer (a.out):

```
$ echo "hello world!" | ./a.out
lexeme:  hello, len=5
lexeme:  world, len=5
!
```

How is the Specification File Converted to a Lexer

□ The problem we face is

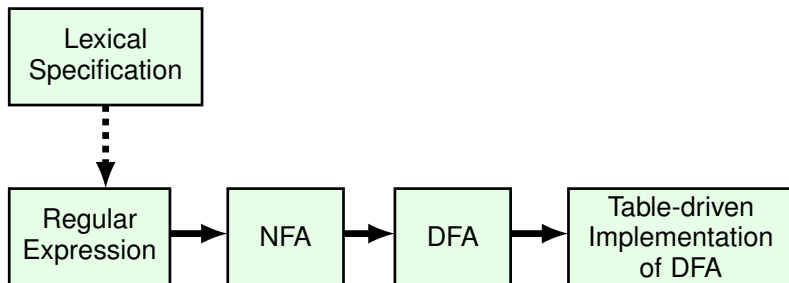
Given a string **s** and a regular expression **RE**, is

$$\mathbf{s} \in \mathbf{L(RE)} ?$$

Implementing Lexical Analysis with Finite Automata

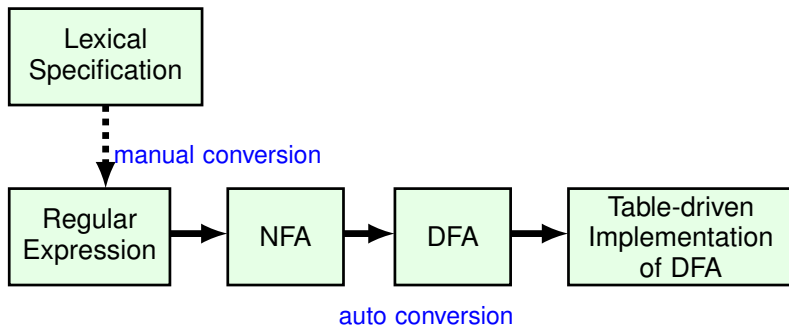
An Overview of RE to FA

Our implementation sketch



An Overview of RE to FA

Our implementation sketch



Implementation Outline

□ RE \Rightarrow NFA \Rightarrow DFA \Rightarrow Table-driven Implementation

- Tokens are specified using Regular Expressions
- Tokens are accepted using a table-driven DFA
 - Deterministic Finite Automata (DFAs)
 - Non-deterministic Finite Automata (NFAs)
- Table implementations

□ I will soon show $RL \equiv L(RE) \equiv L(NFA) \equiv L(DFA)$

- Will show an automated way to take RE all the way to DFA

Finite Automata

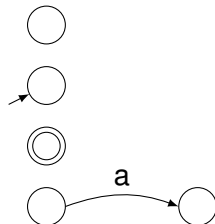
□ A finite automata consists of 5 components
 $(\Sigma, S, n, F, \delta)$

- (1). An input alphabet Σ
- (2). A set of states S
- (3). A start state $n \in S$
- (4). A set of accepting states $F \subseteq S$
- (5). A set of transitions $\delta: S_a \xrightarrow{\text{input}} S_b$

State Graph is a graphical representation of FA

- Sometimes we use **state graph** to represent a FA
- A **state graph** includes

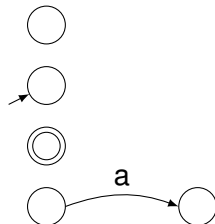
- A set of states
- A start state
- A set of accepting states
- A set of transitions



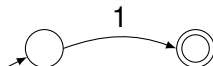
State Graph is a graphical representation of FA

- Sometimes we use **state graph** to represent a FA
- A **state graph** includes

- A set of states
- A start state
- A set of accepting states
- A set of transitions

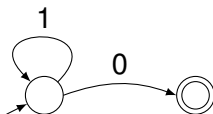


- Example: a finite state automata that accepts only "1"




More State Graph Examples

- A finite automata accepting any number of **1**s followed by a single **0**. Here we have Alphabet = $\{0,1\}$



More State Graph Examples


- A finite automata accepting any number of **1**s followed by a single **0**. Here we have Alphabet = $\{0,1\}$




- Example: What language does the following state graph recognize? Here we have Alphabet = $\{0,1\}$

More State Graph Examples

- A finite automata accepting any number of **1**s followed by a single **0**. Here we have Alphabet = $\{0,1\}$



- Example: What language does the following state graph recognize? Here we have Alphabet = $\{0,1\}$



How tokens are accepted by Finite Automata

Transition $\delta: S_a \xrightarrow{c} S_b$ read as:

When in state S_a , go to state S_b when scanned char is “c”

- 1 Begin at start state $n \in S$.
- 2 Transition until end of input or no transition possible
- 3 Read current state x
 - If $x \in$ accepting set F , then \Rightarrow **accept**
 - Otherwise, \Rightarrow **reject**

Table Implementation of a DFA

Given the state graph of a DFA,

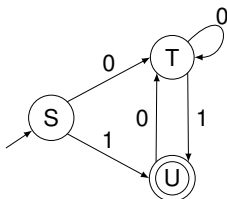
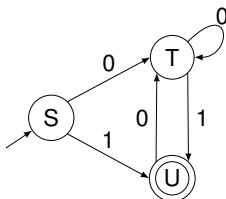


Table Implementation of a DFA

Given the state graph of a DFA,




→ input characters

state ↓

	0	1
S		
T		
U		

Table Implementation of a DFA

Given the state graph of a DFA,



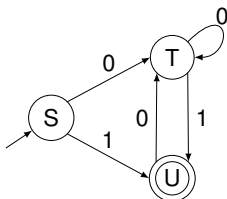
→ input characters

state ↓

	0	1
S	T	U
T	T	U
U	T	X

Table Implementation of a DFA

Given the state graph of a DFA,



→ input characters

state ↓

	0	1
S	T	U
T	T	U
U	T	X

Table-driven Code:

```

DFA() {
    state = "S";
    while (!done) {
        ch = fetch_input();
        state = Table[state][ch];
        if (state == "x")
            perror("error");
    }
    if (state ∈ F)
        printf("accept");
    else
        printf("reject");
}
  
```

Table-driven Code is identical across DFAs

- ❑ Each RE has a different DFA, meaning different table
- ❑ But table-driven code for tokenization remains the same!

Table-driven Code is identical across DFAs

- Each RE has a different DFA, meaning different table

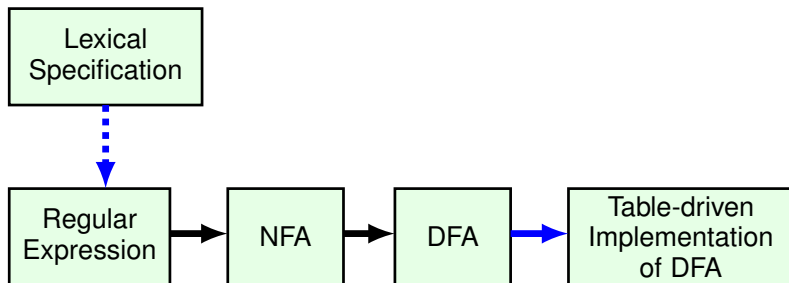
- But table-driven code for tokenization remains the same!

- Revisit our implementation outline

RE \Rightarrow NFA \Rightarrow **DFA** \Rightarrow **Table-driven Implementation**

From RE to FA


Our implementation sketch



Epsilon Moves

■ Another kind of transition: ϵ -moves

- Moves from state A to state B without reading input




■ ϵ -moves makes transitions nondeterministic

- Instead of next character determining the next state, now machine has a choice of staying in A or moving to B

Deterministic and Nondeterministic Automata

- ❑ Deterministic Finite Automata (DFA)
 - One transition per input per state
 - No ϵ -moves
- ❑ Non-deterministic Finite Automata (NFA)
 - Can have multiple transitions for one input in a given state
 - Can have ϵ -moves
- ❑ DFAs are easier to implement
 - Real machines are deterministic by nature just like DFAs
 - Given a nondeterministic choice, hard to know what to do

Examples




Converting RE to NFA


McNaughton-Yamada-Thompson Algorithm

Step 1: processing atomic REs

➤ ϵ expression



➤ single character RE a



Converting RE to NFA (cont.)

Step 2: processing compound REs

➤ $r = s \mid t$


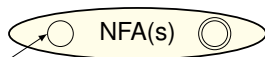
➤ $r = s t$

➤ $r = s^*$

Converting RE to NFA (cont.)

Step 2: processing compound REs

➤ $r = s \mid t$



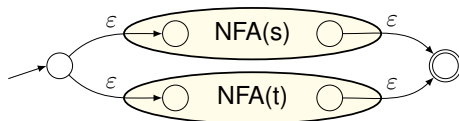
➤ $r = s t$

➤ $r = s^*$

Converting RE to NFA (cont.)

Step 2: processing compound REs

➤ $r = s \mid t$



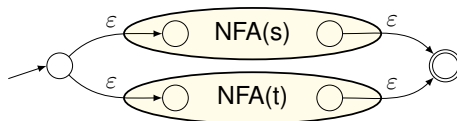
➤ $r = s t$

➤ $r = s^*$

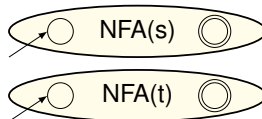
Converting RE to NFA (cont.)

Step 2: processing compound REs

➤ $r = s \mid t$



➤ $r = s t$




➤ $r = s^*$

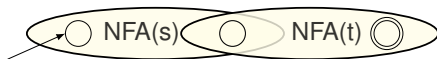
Converting RE to NFA (cont.)

Step 2: processing compound REs

➤ $r = s \mid t$



➤ $r = s t$




➤ $r = s^*$


Converting RE to NFA (cont.)

Step 2: processing compound REs


➤ $r = s \mid t$



➤ $r = s t$




➤ $r = s^*$




Converting RE to NFA (cont.)

Step 2: processing compound REs

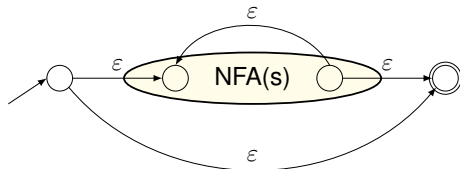
➤ $r = s \mid t$




➤ $r = s t$



➤ $r = s^*$




In-class Practice

 Convert “**(a|b)* a b b**” to NFA

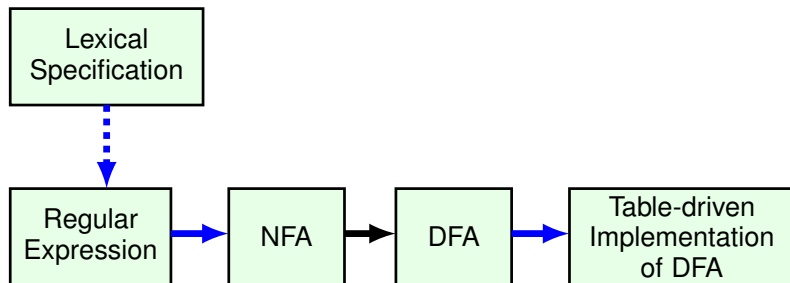
In-class Practice

Convert “**(a|b)* a b b**” to NFA



From RE to FA

Our implementation sketch



Execution of Finite Automata

- ❑ A DFA can take only one path through the state graph
 - Completely determined by input
- ❑ A NFA has a choice of:
 - Whether to make ϵ -moves
 - Which of multiple transitions for a single input to take

Execution of Finite Automata



- ❑ A DFA can take only one path through the state graph
 - Completely determined by input
- ❑ A NFA has a choice of:
 - Whether to make ϵ -moves
 - Which of multiple transitions for a single input to take
- ❑ When to accept a token for DFAs and NFAs
 - DFA: When input string leads to a final state
 - NFA: When input string ****can**** lead to a final state

Execution of Finite Automata

- ❑ A DFA can take only one path through the state graph
 - Completely determined by input
- ❑ A NFA has a choice of:
 - Whether to make ϵ -moves
 - Which of multiple transitions for a single input to take
- ❑ When to accept a token for DFAs and NFAs
 - DFA: When input string leads to a final state
 - NFA: When input string ****can**** lead to a final state
- ❑ **Question:** Which one is more expressive?

$$L(NFA) \equiv L(DFA) \equiv RL$$

Both accept “**(a|b)* a b b**”



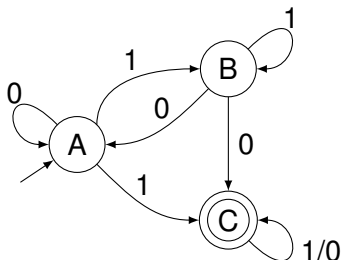
How to Convert NFA to DFA

- ❑ Basic idea: Given NFA, simulate its execution using DFA
 - At step n , the NFA may be in any of multiple possible states
- ❑ The new DFA is constructed as follows,
 - A state of DFA \equiv a non-empty subset of states of the NFA
 - Start state \equiv the set of NFA states reachable through ϵ -moves from NFA start state
 - A transition $S_a \xrightarrow{c} S_b$ is added **iff**

S_b is the set of NFA states reachable from any state in S_a after seeing the input c , considering ϵ -moves as well

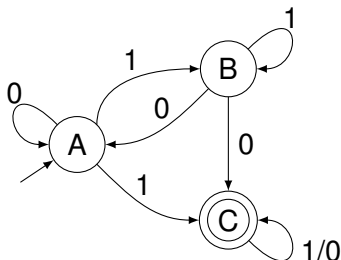
Example NFA to DFA

What is the Equivalent DFA ?



Example NFA to DFA

What is the Equivalent DFA ?

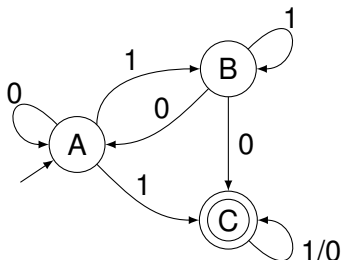


state ↓ → input characters

	0	1
A		
B		
C		

Example NFA to DFA

What is the Equivalent DFA ?

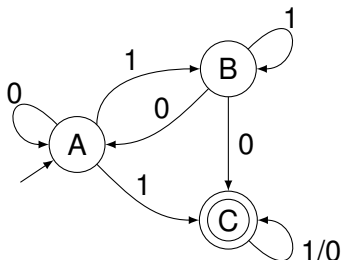


state ↓ → input characters

	0	1
A	A	BC
B	AC	B
C	C	C

Example NFA to DFA

What is the Equivalent DFA ?

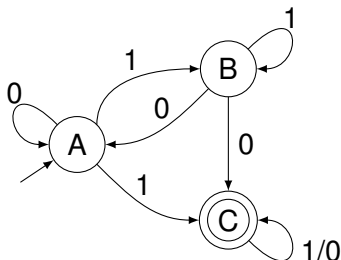


state ↓ → input characters

	0	1
A	A	BC
B	AC	B
C	C	C

Example NFA to DFA

What is the Equivalent DFA ?

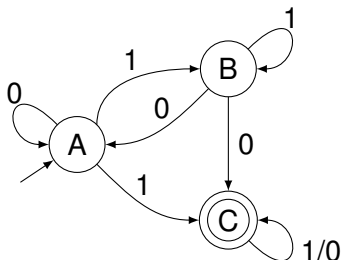


state ↓ → input characters

	0	1
A	A	BC
B	AC	B
C	C	C
AC		
BC		

Example NFA to DFA

What is the Equivalent DFA ?

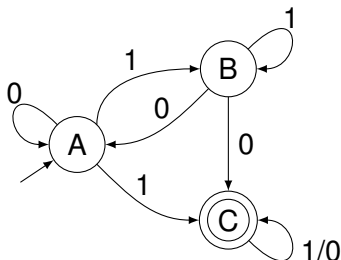


state ↓ → input characters

	0	1
A	A	BC
B	AC	B
C	C	C
AC	AC	BC
BC	AC	BC

Example NFA to DFA

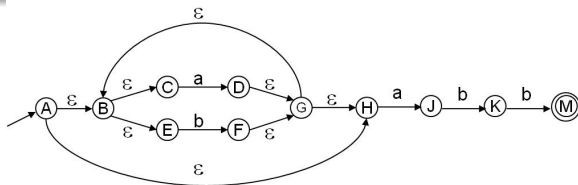
What is the Equivalent DFA ?



state ↓ → input characters

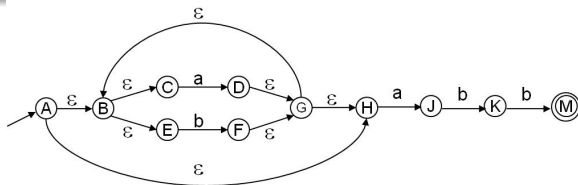
	0	1
A	A	BC
B	AC	B
C	C	C
AC	AC	BC
BC	AC	BC
AB	x	x
ABC	x	x

Algorithm Illustrated: Converting NFA to DFA



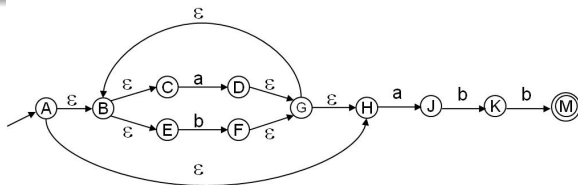
	ϵ	a	b
A			
B			
C			
D			
E			
F			
G			
H			
J			
K			
M			

Step 1: Construct the Table



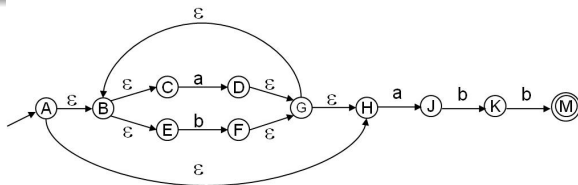
	ϵ	a	b
A	BH		
B	CE		
C		D	
D	G		
E			F
F	G		
G	BH		
H		J	
J			K
K			M
M			

Step 2: Construct ϵ -closure



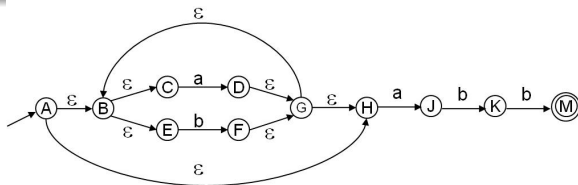
	ϵ	a	b
A	BHCE		
B	CE		
C		D	
D	GBHCE		
E			F
F	GBHCE		
G	BHCE		
H		J	
J			K
K			M
M			

Step 3: Update Other Columns



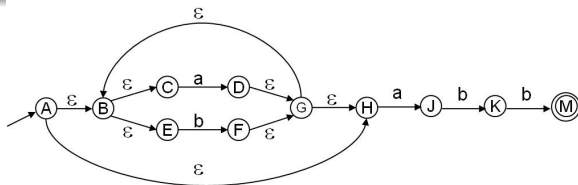
	ϵ	a	b
A	BHCE	DJ	F
B	CE	D	F
C		D	
D	GBHCE	DJ	F
E			F
F	GBHCE	DJ	F
G	BHCE	DJ	F
H		J	
J			K
K			M
M			

Step 4: Construct a New Table



	ϵ	a	b
A	BHCE	DJ	F
B	CE	D	F
C		D	
D	GBHCE	DJ	F
E			F
F	GBHCE	DJ	F
G	BHCE	DJ	F
H		J	
J			K
K			M
M			

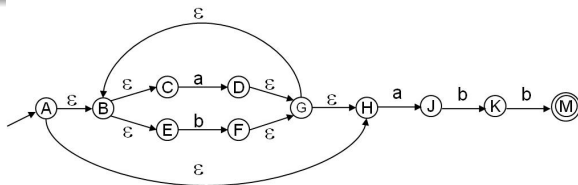
Step 4: Construct a New Table



	ϵ	a	b
A	BHCE	DJ	F
B	CE	D	F
C		D	
D	GBHCE	DJ	F
E			F
F	GBHCE	DJ	F
G	BHCE	DJ	F
H		J	
J			K
K			M
M			

	a	b
ABHCE	DJ	F

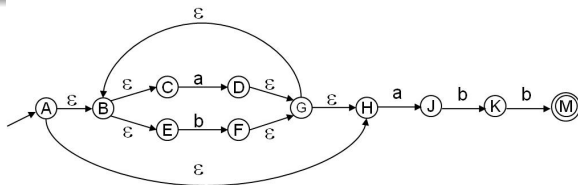
Step 4: Construct a New Table



	ϵ	a	b
A	BHCE	DJ	F
B	CE	D	F
C		D	
D	GBHCE	DJ	F
E			F
F	GBHCE	DJ	F
G	BHCE	DJ	F
H		J	
J			K
K			M
M			

	a	b
ABHCE	DJ	F
DJ	DJ	FK
F	DJ	F

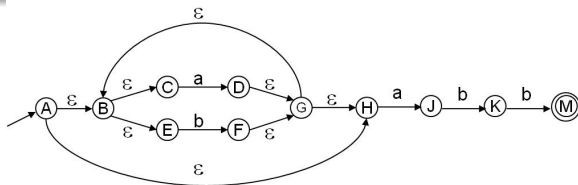
Step 4: Construct a New Table



	ϵ	a	b
A	BHCE	DJ	F
B	CE	D	F
C		D	
D	GBHCE	DJ	F
E			F
F	GBHCE	DJ	F
G	BHCE	DJ	F
H		J	
J			K
K			M
M			

	a	b
ABHCE	DJ	F
DJ	DJ	FK
F	DJ	F
FK	DJ	FM

Step 4: Construct a New Table



	ϵ	a	b
A	BHCE	DJ	F
B	CE	D	F
C		D	
D	GBHCE	DJ	F
E			F
F	GBHCE	DJ	F
G	BHCE	DJ	F
H		J	
J			K
K			M
M			


	a	b
ABHCE	DJ	F
DJ	DJ	FK
F	DJ	F
FK	DJ	FM
FM	DJ	F


Step 5: Generate the DFA

	a	b
ABHCE	DJ	F
DJ	DJ	FK
F	DJ	F
FK	DJ	FM
FM	DJ	F

Step 5: Generate the DFA

	a	b
ABHCE	DJ	F
DJ	DJ	FK
F	DJ	F
FK	DJ	FM
FM	DJ	F


 Note: the number of states is not minimized



Step 5: Generate the DFA

	a	b
ABHCE	DJ	F
DJ	DJ	FK
F	DJ	F
FK	DJ	FM
FM	DJ	F

□ Note: the number of states is not minimized



□ States ABHCE and F can be merged! Why?

NFA to DFA. Space Complexity

- ❑ An NFA may be in many states at any time
- ❑ If NFA has N states, how many DFA states?
 - NFA must be in some subset of those N states
 - How many non-empty subsets are there ?
 - $2^N - 1$ many states
- ❑ The resulting DFA has $O(2^N)$ space complexity
 - Remember this is big-O. Typically much less.


NFA to DFA Time Complexity

- ❑ A DFA can be implemented by a 2D table T
 - One dimension is “states”, the other is “input characters”
 - For $S_a \xrightarrow{c} S_b$, we have $T[S_a, c] = S_b$
- ❑ DFA execution
 - If the current state is S_a and input is c , then read $T[S_a, c]$
 - Update the current state to S_b , assuming $T[S_a, c] = S_b$
 - Requires $O(|X|)$ steps, where $|X|$ is the length of input
- ❑ NFA execution
 - At a given step, there is a set of possible states, up to N
 - On input c , must access table for each possible state to get next set of possible states
 - Requires $O(|X| * N)$ steps

Implementation in Practice

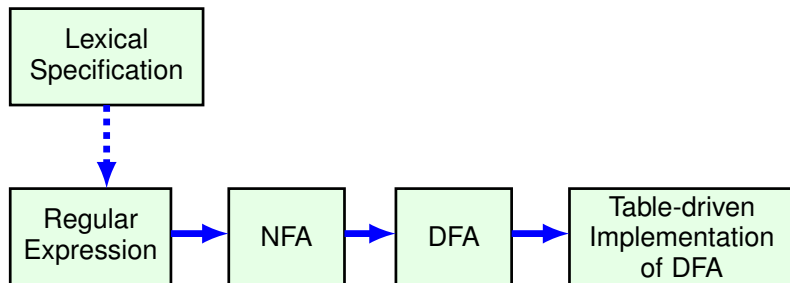
GNU **lex**

- Convert regular expression to NFA
- Convert NFA to DFA
- Perform DFA state minimization to reduce space
- Generate transition table from DFA
- Perform table compression to further reduce space

 Most other automated lexers also trade off space for speed by choosing DFA over NFA

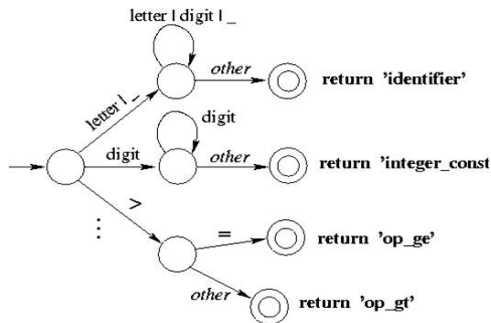
From RE to FA

Our implementation sketch



Structure of a Scanner Automaton

■ A scanner recognize multiple REs



How much should we match?

- ❑ In general, find the longest match possible
- ❑ If same length, first rule in lex file takes precedence

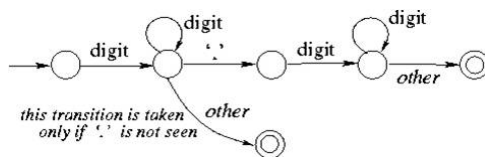
Example:

on input **123.45**, we match it as

(numConst, 123.45)

rather than

(numConst, 123), (dot, "."), (numConst, 45)



How to Match Keywords?

- ❑ Approach 1: Hardcode the keywords
- ❑ Approach 2: When token is identified, check a table of keywords

Example: to recognize the following tokens

Identifiers: `letter(letter|digit)*`

Keywords: `if, then, else`

Beyond Regular Languages

- ❑ Regular languages are expressive enough for tokens
- ❑ Can express identifiers, strings, comments, etc.
- ❑ However, it is the least expressive formal language .
 - Many languages are not regular
 - C programming language is not
 - XML or JSON is not
 - “(((...)))” is also not
 - But YAML is!
 - Finite automata cannot express nested structures
 - Can only remember a finite number of “(”s encountered
- ❑ We need a more powerful language for parsing
 - In the next lecture, we will discuss **context-free languages**